

# Assignment 1

Nicholas Jacob

2024-08-19

```
knitr::opts_chunk$set(echo = TRUE)
library(plyr)
```

## Using R: Vectors

Using `c` to combine the values, we see that  $x$  is a vector.

```
x<- c(3,12,6,-5,0,8,15,1,-10,7)
is.vector(x)
```

```
## [1] TRUE
```

To create the new vector  $y$  as a sequence from the min of  $x$  to the max of  $x$ , we do the following:

```
y <-seq(min(x),max(x), length.out = 10)
y
```

```
## [1] -10.000000 -7.222222 -4.444444 -1.666667 1.111111 3.888889
## [7] 6.666667 9.444444 12.222222 15.000000
```

I was not familiar with the `length.out` command but found it in the `Help` package to see that it would restrict the output to that many elements.

We compute the desired stats next

```
#consider changing this one with some tidy code
sum(x)
```

```
## [1] 37
```

```
sum(y)
```

```
## [1] 25
```

```
mean(x)
```

```
## [1] 3.7
```

```
mean(y)
```

```
## [1] 2.5
```

```
sd(x)
```

```
## [1] 7.572611
```

```
sd(y)
```

```
## [1] 8.41014
```

```
var(x)
```

```
## [1] 57.34444
```

```
var(y)
```

```
## [1] 70.73045
```

```
mad(x)
```

```
## [1] 5.9304
```

```
mad(y)
```

```
## [1] 10.29583
```

```
quantile(x,1/4)
```

```
## 25%
```

```
## 0.25
```

```
quantile(y,1/4)
```

```
## 25%
```

```
## -3.75
```

```
quantile(x,3/4)
```

```
## 75%
```

```
## 7.75
```

```
quantile(y,3/4)
```

```
## 75%
```

```
## 8.75
```

```
quantile(x,1/5)
```

```
## 20%
```

```
## -1
```

```
quantile(y,1/5)
```

```
## 20%
```

```
## -5
```

```
quantile(x,3/5)
```

```
## 60%
```

```
## 6.4
```

```
quantile(y,3/5)
```

```
## 60%
```

```
## 5
```

```
quantile(x,2/5)
```

```
## 40%
```

```
## 2.2
```

```
quantile(y,2/5)
```

```
##          40%
## -1.665335e-15
```

```
quantile(x,4/5)
```

```
## 80%
## 8.8
```

```
quantile(y,4/5)
```

```
## 80%
## 10
```

To do sampling with replacement we do the following

```
sample(x,7,TRUE)
```

```
## [1]  7  3  7 -5  1  8  0
```

The TRUE gives the replacement. Some instances do see repeated vales.

Next we do the `t.test`

```
t.test(x,y)
```

```
##
## Welch Two Sample t-test
##
## data:  x and y
## t = 0.33531, df = 17.805, p-value = 0.7413
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -6.324578  8.724578
## sample estimates:
## mean of x mean of y
##      3.7      2.5
```

We fail to reject the null hypothesis here. There is no evidence to suggest that the mean values are different.

Next we explore the `order` function.

```
order(x)
```

```
## [1]  9  4  5  8  1  3 10  6  2  7
```

We see this gives the order of the elements of  $x$ , indexing at 1 as the lowest value. To sort  $x$  we could do the following.

```
sort(x)
```

```
## [1] -10 -5  0  1  3  6  7  8 12 15
```

We could also use the order function as follows:

```
x[order(x)]
```

```
## [1] -10 -5  0  1  3  6  7  8 12 15
```

Inside the `[]` we are giving the index of the value we want. So this will return the values in the proper order. Lastly we will preform the paired t.test.

```
t.test(sort(x),y,paired = TRUE)
```

```
##
```

```
## Paired t-test
##
## data: sort(x) and y
## t = 2.164, df = 9, p-value = 0.05868
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
## -0.05440584 2.45440584
## sample estimates:
## mean difference
## 1.2
```

The result here is still not significant (for  $p = 0.05$ ) but is much closer than in the non-paired data. I am actually quite surprised at that result but since  $y$  is build off of  $x$  and now they are both sequential I could see why they might be statistically equivalent on average.

A logical test for negativity is simply

```
x>0
```

```
## [1] TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE FALSE TRUE
```

Since this gives the Boolean, we can use that as the index for  $x$  and overwrite  $x$

```
x <- x[x>0]
```

```
x
```

```
## [1] 3 12 6 8 15 1 7
```

## Using R: Some Missing Values

```
col1 <- c(1,2,3,NA,5)
col2 <- c(4,5,6,89,101)
col3 <- c(45,NA,66,121,201)
col4 <- c(14,NA,13,NA,27)
X <- rbind (col1,col2,col3,col4)
```

```
X
```

```
##      [,1] [,2] [,3] [,4] [,5]
## col1  1   2   3  NA   5
## col2  4   5   6  89  101
## col3  45  NA  66  121  201
## col4  14  NA  13  NA   27
```

So we see  $X$  has NA in three rows. We can find the NAs with the following

```
is.na(X)
```

```
##      [,1] [,2] [,3] [,4] [,5]
## col1 FALSE FALSE FALSE TRUE FALSE
## col2 FALSE FALSE FALSE FALSE FALSE
## col3 FALSE TRUE FALSE FALSE FALSE
## col4 FALSE TRUE FALSE TRUE FALSE
```

To get to which rows have the NAs, we sum across the booleans and ask that the sum in that row is larger than 0. Then we use the rownames command to give out those rows names that do have some NAs.

```
rownames(X)[rowSums(is.na(X))>0]
```

```
## [1] "col1" "col3" "col4"
```

For the next piece, we define  $y$

```
y <- c(3,12,99,99,7,99,21)
y
```

```
## [1] 3 12 99 99 7 99 21
```

We will find the 99s with this peice of code

```
y == 99
```

```
## [1] FALSE FALSE TRUE TRUE FALSE TRUE FALSE
```

We set that to the NA value with this which overwrites y values.

```
y[y==99] = NA
```

```
y
```

```
## [1] 3 12 NA NA 7 NA 21
```

I count the NA values with a sum of the booleans

```
sum(is.na(y))
```

```
## [1] 3
```

##Using R: IDE

Here I have read the data in. I utilize the head command to display the first 6 rows.

```
college = read.csv('college.csv')
head(college)
```

```
##               X Private Apps Accept Enroll Top10perc Top25perc
## 1 Abilene Christian University    Yes 1660   1232    721      23      52
## 2      Adelphi University        Yes 2186   1924    512      16      29
## 3      Adrian College          Yes 1428   1097    336      22      50
## 4      Agnes Scott College       Yes  417    349    137      60      89
## 5 Alaska Pacific University     Yes  193    146     55      16      44
## 6      Albertson College        Yes  587    479    158      38      62
##  F.Undergrad P.Undergrad Outstate Room.Board Books Personal PhD Terminal
## 1      2885      537    7440    3300    450    2200    70      78
## 2      2683    1227    12280    6450    750    1500    29      30
## 3      1036      99    11250    3750    400    1165    53      66
## 4      510      63    12960    5450    450    875    92      97
## 5      249    869    7560    4120    800    1500    76      72
## 6      678      41    13500    3335    500    675    67      73
##  S.F.Ratio perc.alumni Expend Grad.Rate
## 1      18.1      12    7041      60
## 2      12.2      16   10527      56
## 3      12.9      30    8735      54
## 4      7.7      37   19016      59
## 5     11.9       2   10922      15
## 6      9.4      11    9727      55
```

Next, I change the rownames to the university name and delete that column.

```
rownames(college) <- college[,1]
college <- college[,-1]
head(college)
```

##	Private	Apps	Accept	Enroll	Top10perc	Top25perc
## Abilene Christian University	Yes	1660	1232	721	23	52
## Adelphi University	Yes	2186	1924	512	16	29
## Adrian College	Yes	1428	1097	336	22	50
## Agnes Scott College	Yes	417	349	137	60	89
## Alaska Pacific University	Yes	193	146	55	16	44
## Albertson College	Yes	587	479	158	38	62

##	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books
## Abilene Christian University	2885	537	7440	3300	450
## Adelphi University	2683	1227	12280	6450	750
## Adrian College	1036	99	11250	3750	400
## Agnes Scott College	510	63	12960	5450	450
## Alaska Pacific University	249	869	7560	4120	800
## Albertson College	678	41	13500	3335	500

##	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend
## Abilene Christian University	2200	70	78	18.1	12	7041
## Adelphi University	1500	29	30	12.2	16	10527
## Adrian College	1165	53	66	12.9	30	8735
## Agnes Scott College	875	92	97	7.7	37	19016
## Alaska Pacific University	1500	76	72	11.9	2	10922
## Albertson College	675	67	73	9.4	11	9727

##	Grad.Rate
## Abilene Christian University	60
## Adelphi University	56
## Adrian College	54
## Agnes Scott College	59
## Alaska Pacific University	15
## Albertson College	55

Next we examine some stats on the data

```
summary(college)
```

##	Private	Apps	Accept	Enroll
## Length:777	Min. : 81	Min. : 72	Min. : 35	
## Class :character	1st Qu.: 776	1st Qu.: 604	1st Qu.: 242	
## Mode :character	Median : 1558	Median : 1110	Median : 434	
##	Mean : 3002	Mean : 2019	Mean : 780	
##	3rd Qu.: 3624	3rd Qu.: 2424	3rd Qu.: 902	
##	Max. : 48094	Max. : 26330	Max. : 6392	

##	Top10perc	Top25perc	F.Undergrad	P.Undergrad
## Min. : 1.00	Min. : 9.0	Min. : 139	Min. : 1.0	
## 1st Qu.:15.00	1st Qu.: 41.0	1st Qu.: 992	1st Qu.: 95.0	
## Median :23.00	Median : 54.0	Median : 1707	Median : 353.0	
## Mean :27.56	Mean : 55.8	Mean : 3700	Mean : 855.3	
## 3rd Qu.:35.00	3rd Qu.: 69.0	3rd Qu.: 4005	3rd Qu.: 967.0	
## Max. :96.00	Max. :100.0	Max. :31643	Max. :21836.0	

##	Outstate	Room.Board	Books	Personal
## Min. : 2340	Min. :1780	Min. : 96.0	Min. : 250	
## 1st Qu.: 7320	1st Qu.:3597	1st Qu.: 470.0	1st Qu.: 850	
## Median : 9990	Median :4200	Median : 500.0	Median :1200	
## Mean :10441	Mean :4358	Mean : 549.4	Mean :1341	
## 3rd Qu.:12925	3rd Qu.:5050	3rd Qu.: 600.0	3rd Qu.:1700	
## Max. :21700	Max. :8124	Max. :2340.0	Max. :6800	

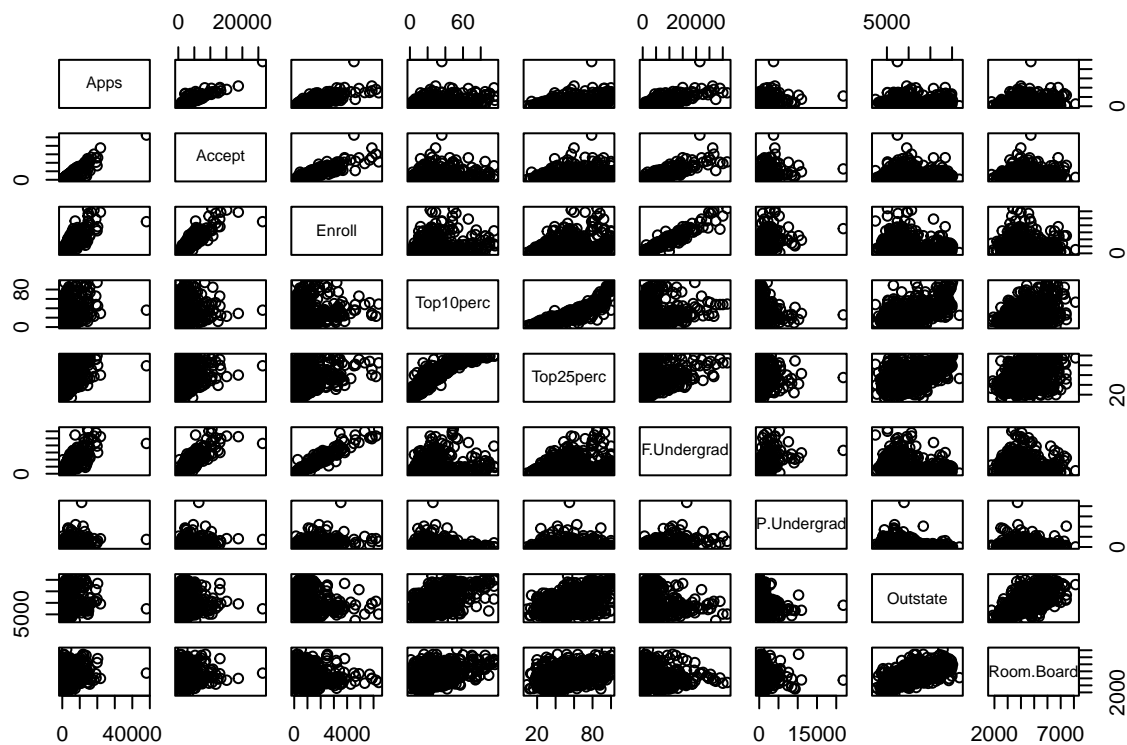
  

##	PhD	Terminal	S.F.Ratio	perc.alumni
##				

```
## Min.    : 8.00    Min.    : 24.0    Min.    : 2.50    Min.    : 0.00
## 1st Qu.: 62.00    1st Qu.: 71.0    1st Qu.:11.50    1st Qu.:13.00
## Median : 75.00    Median : 82.0    Median :13.60    Median :21.00
## Mean   : 72.66    Mean   : 79.7    Mean   :14.09    Mean   :22.74
## 3rd Qu.: 85.00    3rd Qu.: 92.0    3rd Qu.:16.50    3rd Qu.:31.00
## Max.   :103.00    Max.   :100.0    Max.   :39.80    Max.   :64.00
##      Expend      Grad.Rate
## Min.    : 3186    Min.    : 10.00
## 1st Qu.: 6751    1st Qu.: 53.00
## Median : 8377    Median : 65.00
## Mean   : 9660    Mean   : 65.46
## 3rd Qu.:10830    3rd Qu.: 78.00
## Max.   :56233    Max.   :118.00
```

I am not familiar with the `pairs` command but here goes

```
pairs(college[,2:10])
```




That is a nice graphic although a bit too small for my tastes. I hope it compiles correctly in the pdf...

Next I'll create the boxplot for out of state tuition vs the public or private.

```
boxplot(Outstate ~ Private, data = college, main = "Out Of State Tuition by College Type", ylab = "Out of State Tuition")
```

## Out Of State Tuition by College Type



This looks fine although I do prefer ggplot2.

Next I comment the code as requested

```
Elite <- rep ("No", nrow(college )) #This creates a vector that full of No that is the same width as th
Elite [college$Top10perc >50] <- "Yes" #this changes some of the nos to yes if the top10 is more than 5
Elite <- as.factor (Elite) #this casts the vector as a factor vector. This is useful in that Elite now
college <- data.frame(college ,Elite) #this adds the column to the original dataframe and saves it

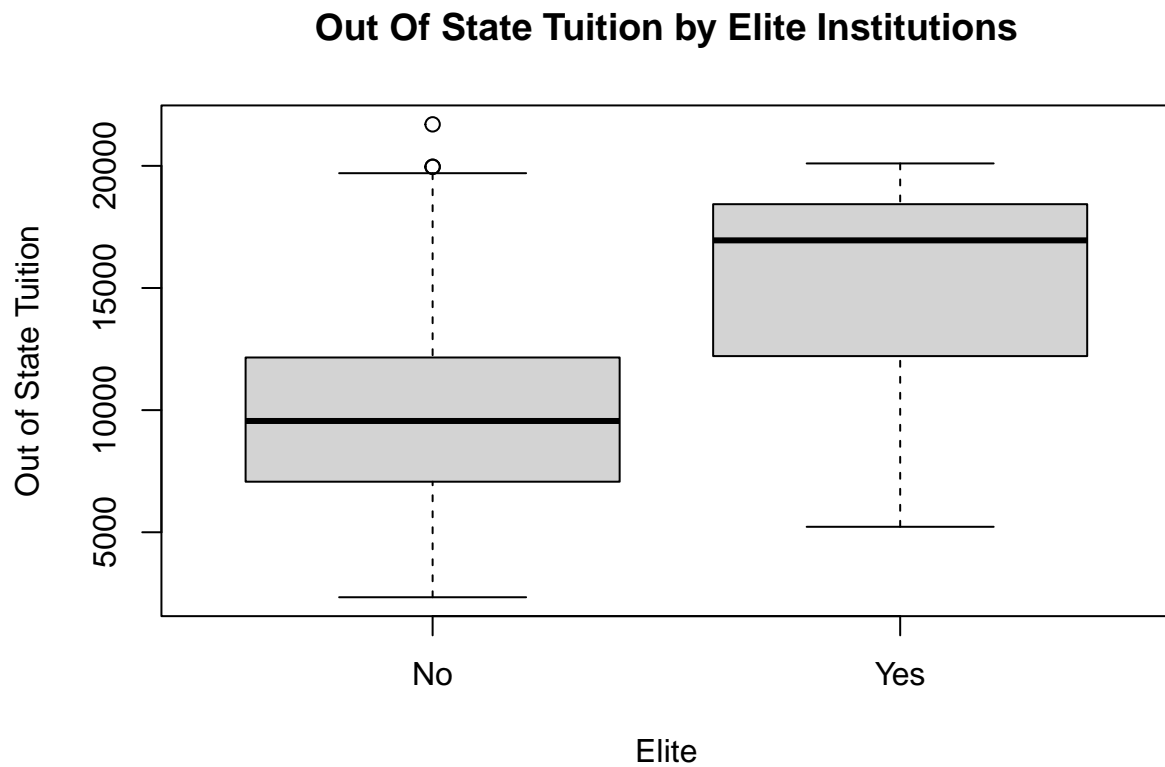
summary(Elite)
```

```
## No Yes
## 699 78
```

It appears that there are 78 elite universities. Let's explore tutions with this new factor

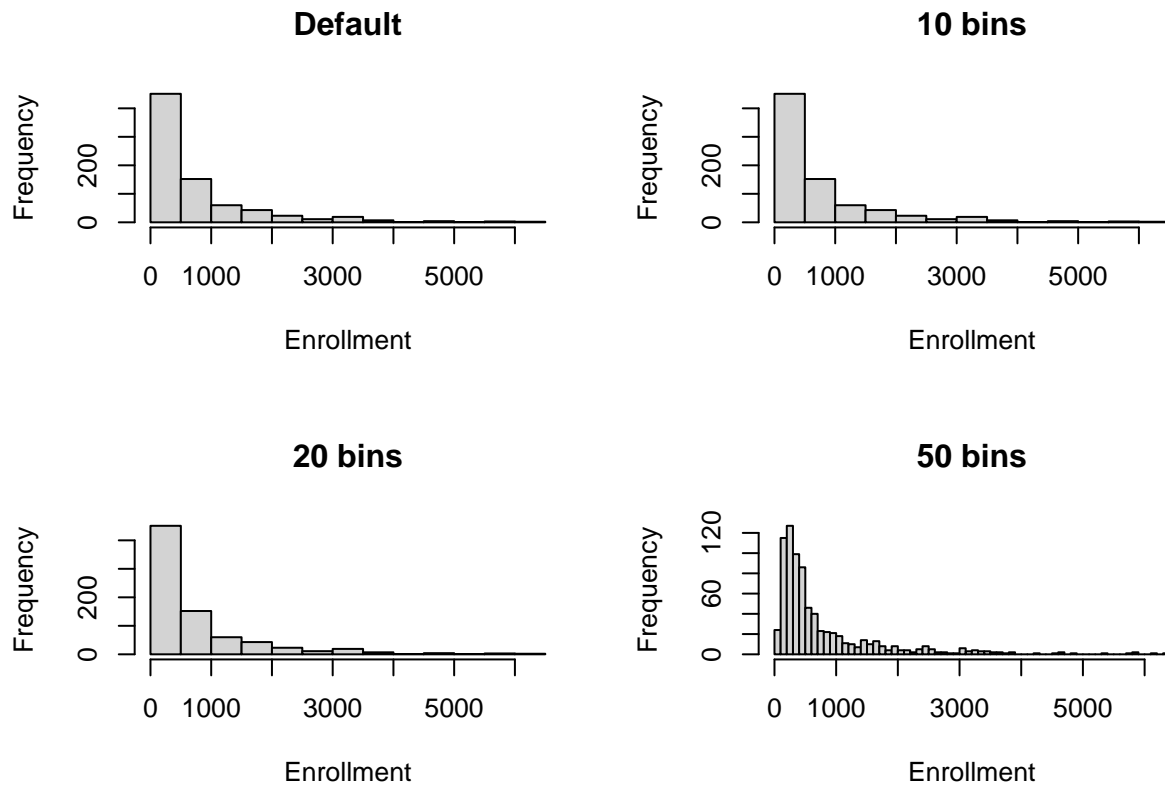
```
boxplot(Outstate ~ Elite, data = college, main = "Out Of State Tuition by Elite Institutions", ylab = "Out of State Tuition")
```





Next we look at a few histograms with differing number of bins.

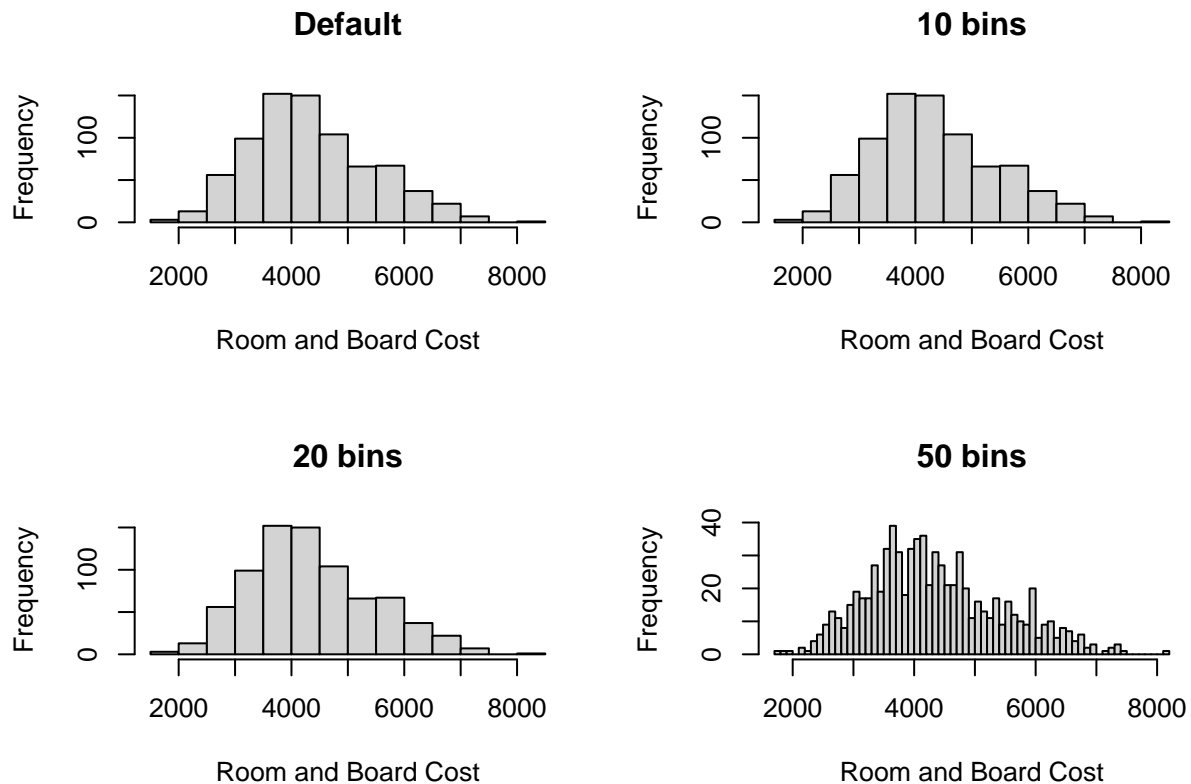
```
par(mfrow=c(2,2))
hist(college[, 'Enroll'], main = "Default", xlab = "Enrollment")
hist(college[, 'Enroll'], main = "10 bins", breaks = 10, xlab = "Enrollment")
hist(college[, 'Enroll'], main = "20 bins", breaks = 20, xlab = "Enrollment")
hist(college[, 'Enroll'], main = "50 bins", breaks = 50, xlab = "Enrollment")
```



I don't see much difference between the default, 10 nor 20. The 50 does look a bit different.

Again just to try it once more

```
par(mfrow=c(2,2))
hist(college[, 'Room.Board'], main = "Default", xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "10 bins", breaks = 10, xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "20 bins", breaks = 20, xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "50 bins", breaks = 50, xlab = "Room and Board Cost")
```



It kind of looks like more than 10 breaks Maybe the default overrides that option if you set it too low...

## Using R: Manipulating Data in Data Frames

First, I'll load some data directly from a package. This `baseball` data comes from the `plyr` package loaded earlier.

```
head(baseball)
```

```
##           id year stint team lg  g  ab  r  h X2b X3b hr rbi sb cs bb so ibb
## 4   ansonca01 1871     1  RC1   25 120 29 39  11   3  0  16  6  2  2  1  NA
## 44  forceda01 1871     1  WS3   32 162 45 45   9   4  0  29  8  0  4  0  NA
## 68  mathebo01 1871     1  FW1   19  89 15 24   3   1  0  10  2  1  2  0  NA
## 99  startjo01 1871     1  NY2   33 161 35 58   5   1  1  34  4  2  3  0  NA
## 102 suttoez01 1871     1  CL1   29 128 35 45   3   7  3  23  3  1  1  0  NA
## 106 whitede01 1871     1  CL1   29 146 40 47   6   5  1  21  2  2  4  1  NA
##      hbp sh sf gidp
## 4      NA NA NA   NA
## 44     NA NA NA   NA
## 68     NA NA NA   NA
## 99     NA NA NA   NA
## 102    NA NA NA   NA
## 106    NA NA NA   NA
```

Lots of baseball data!

```
baseball[baseball$year<1954,'sf'] = 0 #set all sf before 1954 to 0
baseball[is.na(baseball$hbp),'hbp'] = 0 #set all null values for hit by pitch to 0
```

```
baseball <- baseball[baseball$ab>=50,]
```

Now that the data is clean, we will apply the obp formula of

$$obp = \frac{h + bb + hbp}{ab + bb + hbp + sf}$$

```
baseball <- mutate(baseball, obp = (h+bb+hbp)/(ab+bb+hbp+sf))
```

```
head(baseball)
```

```
##           id year stint team lg  g  ab  r  h X2b X3b hr rbi sb cs bb so ibb
## 4  ansonca01 1871     1  RC1   25 120 29 39  11   3  0  16  6  2  2  1  NA
## 44 forceda01 1871     1  WS3   32 162 45 45   9   4  0  29  8  0  4  0  NA
## 68 mathebo01 1871     1  FW1   19  89 15 24   3   1  0  10  2  1  2  0  NA
## 99 startjo01 1871     1  NY2   33 161 35 58   5   1  1  34  4  2  3  0  NA
## 102 suttoez01 1871     1  CL1   29 128 35 45   3   7  3  23  3  1  1  0  NA
## 106 whitede01 1871     1  CL1   29 146 40 47   6   5  1  21  2  2  4  1  NA
##           hbp sh sf gidp      obp
## 4           0 NA  0   NA 0.3360656
## 44          0 NA  0   NA 0.2951807
## 68          0 NA  0   NA 0.2857143
## 99          0 NA  0   NA 0.3719512
## 102         0 NA  0   NA 0.3565891
## 106         0 NA  0   NA 0.3400000
```

Now that we have that info added, let's find the top five players for obp of all time.

```
arrange(baseball, -obp)[1:5,c('year','id','obp')] #I get the top records with 1:5, restrict on to the c
```

```
##   year      id      obp
## 1 2004 bondsba01 0.6094003
## 2 2002 bondsba01 0.5816993
## 3 1941 willite01 0.5528053
## 4 1899 mcgrajo01 0.5474860
## 5 1923 ruthba01 0.5445402
```

We see here Barry Bonds (from the 'roids era twice), Ted Williams(a year he hit .400), John McGraw (a player I was not familiar with though he did have a season with my home team Cardinals in 1900) and the babe himself Babe Ruth.

##Using R: aggregate() Function

I am going to grab the quakes dataset.


```
head(quakes)
```

```
##      lat   long depth mag stations
## 1 -20.42 181.62   562 4.8         41
## 2 -20.62 181.03   650 4.2         15
## 3 -26.00 184.10    42 5.4         43
## 4 -17.97 181.66   626 4.1         19
## 5 -20.42 181.96   649 4.0         11
## 6 -19.68 184.31   195 4.0         12
```

Next we will examine magnitude versus depth with a scatter plot.

```
plot(quakes$depth,quakes$mag, xlab = 'Depth', ylab = 'Magnitude', main = 'Scatter Plot of Depth vs Magn
```

## Scatter Plot of Depth vs Magnitude



Next we will aggregate the data to look at the average depth for each of the magnitude levels

```
quakeAvgDepth = aggregate(quakes$depth, list(mag = quakes$mag), mean)
```

Not too bad when you follow the example in the help menu.

Next I rename the dataframe to have useful column names and print it to see the nice output.


```
colnames(quakeAvgDepth) = c('mag', 'meanDepth')
head(quakeAvgDepth)
```

```
##   mag meanDepth
## 1 4.0  410.0652
## 2 4.1  412.4000
## 3 4.2  389.8778
## 4 4.3  357.9294
## 5 4.4  307.1188
## 6 4.5  333.6729
```

Now we plot again to see if there is a relationship in the aggregate

```
plot(quakeAvgDepth$mag, quakeAvgDepth$meanDepth, xlab = 'Magnitude', ylab = 'Average Depth', main = 'Scat
```

### Scatter of Aggregated Magnitude vs Mean Depth



There clearly appears to be a relationship here. It was not as obvious in the full data case but the relationship appears in the aggregate. I do question a bit of this methodology though. We are aggregating a continuous variable that has been truncated to two decimals. Richter scale (magnitude) is a famous example of a logarithmic scale so small rounding errors are amplified in varying degrees as you increase the scale. While yes, I believe there is a relationship, I'd be worried about generalizing too far based on this data.