

# Assignment 1

Nicholas Jacob

2024-08-19

```
knitr::opts_chunk$set(echo = TRUE)
library(plyr)
```

## 1 Using R: Vectors

a. Using `c` to combine the values, we see that  $x$  is a vector.

```
x<- c(3,12,6,-5,0,8,15,1,-10,7)
is.vector(x)
```

```
## [1] TRUE
```

b. To create the new vector  $y$  as a sequence from the min of  $x$  to the max of  $x$ , we do the following:

```
y <-seq(min(x),max(x), length.out = 10)
y
```

```
## [1] -10.000000 -7.222222 -4.444444 -1.666667 1.111111 3.888889
## [7] 6.666667 9.444444 12.222222 15.000000
```

I was not familiar with the `length.out` command but found it in the `Help` package to see that it would restrict the output to that many elements.

c. We compute the desired stats next

```
#consider changing this one with some tidy code
sum(x)
```

```
## [1] 37
```

```
sum(y)
```

```
## [1] 25
```

```
mean(x)
```

```
## [1] 3.7
```

```
mean(y)
```

```
## [1] 2.5
```

```
sd(x)
```

```
## [1] 7.572611
```

```
sd(y)
```

```
## [1] 8.41014
```

```
var(x)
```

```
## [1] 57.34444
```

```
var(y)
```

```
## [1] 70.73045
```

```
mad(x)
```

```
## [1] 5.9304
```

```
mad(y)
```

```
## [1] 10.29583
```

```
quantile(x,1/4)
```

```
## 25%
```

```
## 0.25
```

```
quantile(y,1/4)
```

```
## 25%
```

```
## -3.75
```

```
quantile(x,3/4)
```

```
## 75%
```

```
## 7.75
```

```
quantile(y,3/4)
```

```
## 75%
```

```
## 8.75
```

```
quantile(x,1/5)
```

```
## 20%
```

```
## -1
```

```
quantile(y,1/5)
```

```
## 20%
```

```
## -5
```

```
quantile(x,3/5)
```

```
## 60%
```

```
## 6.4
```

```
quantile(y,3/5)
```

```
## 60%
```

```
## 5
```

```
quantile(x,2/5)
```

```
## 40%
```

```
## 2.2
```

```
quantile(y,2/5)
```

```
##          40%
## -1.665335e-15
```

```
quantile(x,4/5)
```

```
## 80%
## 8.8
```

```
quantile(y,4/5)
```

```
## 80%
## 10
```

d. To do sampling with replacement we do the following

```
z <- sample(x,7,TRUE)
z
```

```
## [1] 12 -10 -10  6  1  6 15
```

The TRUE gives the replacement. Some instances do see repeated vales.

e. Next we do the `t.test`

```
t.test(x,y)
```

```
##
## Welch Two Sample t-test
##
## data: x and y
## t = 0.33531, df = 17.805, p-value = 0.7413
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -6.324578  8.724578
## sample estimates:
## mean of x mean of y
##      3.7      2.5
```

We fail to reject the null hypothesis here. There is no evidence to suggest that the mean values are different.

f. Next we explore the `order` function.

```
order(x)
```

```
## [1]  9  4  5  8  1  3 10  6  2  7
```

We see this gives the order of the elements of  $x$ , indexing at 1 as the lowest value. To sort  $x$  we could do the following.

```
sort(x)
```

```
## [1] -10 -5  0  1  3  6  7  8 12 15
```

We could also use the order function as follows:

```
x[order(x)]
```

```
## [1] -10 -5  0  1  3  6  7  8 12 15
```

Inside the `[]` we are giving the index of the value we want. So this will return the values in the proper order. Lastly we will preform the paired t.test.

```
t.test(sort(x),y,paired = TRUE)
```

```
##
## Paired t-test
##
## data: sort(x) and y
## t = 2.164, df = 9, p-value = 0.05868
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
## -0.05440584 2.45440584
## sample estimates:
## mean difference
## 1.2
```

The result here is still not significant (for  $p = 0.05$ ) but is much closer than in the non-paired data. I am actually quite surprised at that result but since  $y$  is build off of  $x$  and now they are both sequential I could see why they might be statistically equivalent on average.

g. A logical test for negativity is simply

```
x>0
```

```
## [1] TRUE TRUE TRUE FALSE FALSE TRUE TRUE TRUE FALSE TRUE
```

Since this gives the Boolean, we can use that as the index for  $x$  and overwrite  $x$

```
x <- x[x>0]
x
```

```
## [1] 3 12 6 8 15 1 7
```

## 2 Using R: Some Missing Values

```
col1 <- c(1,2,3,NA,5)
col2 <- c(4,5,6,89,101)
col3 <- c(45,NA,66,121,201)
col4 <- c(14,NA,13,NA,27)
X <- rbind (col1,col2,col3,col4)
```

```
X
```

```
##      [,1] [,2] [,3] [,4] [,5]
## col1    1    2    3  NA    5
## col2    4    5    6  89   101
## col3   45   NA   66  121  201
## col4   14   NA   13   NA   27
```

a. So we see  $X$  has NA in three rows. We can find the NAs with the following

```
is.na(X)
```

```
##      [,1] [,2] [,3] [,4] [,5]
## col1 FALSE FALSE FALSE TRUE FALSE
## col2 FALSE FALSE FALSE FALSE FALSE
## col3 FALSE TRUE FALSE FALSE FALSE
## col4 FALSE TRUE FALSE TRUE FALSE
```

To get to which rows have the NAs, we sum across the booleans and ask that the sum in that row is larger than 0. Then we use the rownames command to give out those rows names that do have some NAs.

```
rownames(X)[rowSums(is.na(X))>0]
```

```
## [1] "col1" "col3" "col4"
```

b. For the next piece, we define  $y$

```
y <- c(3,12,99,99,7,99,21)
y
```

```
## [1] 3 12 99 99 7 99 21
```

We will find the 99s with this peice of code

```
y == 99
```

```
## [1] FALSE FALSE TRUE TRUE FALSE TRUE FALSE
```

We set that to the NA value with this which overwrites  $y$  values.

```
y[y==99] = NA
y
```

```
## [1] 3 12 NA NA 7 NA 21
```

I count the NA values with a sum of the booleans

```
sum(is.na(y))
```

```
## [1] 3
```

### 3 Using R: IDE

a. Here I have read the data in. I utilize the head command to display the first 6 rows.

```
college = read.csv('college.csv')
head(college)
```

```
##               X Private Apps Accept Enroll Top10perc Top25perc
## 1 Abilene Christian University      Yes 1660   1232    721      23      52
## 2           Adelphi University      Yes 2186   1924    512      16      29
## 3           Adrian College      Yes 1428   1097    336      22      50
## 4           Agnes Scott College      Yes  417    349    137      60      89
## 5      Alaska Pacific University      Yes  193    146     55      16      44
## 6           Albertson College      Yes  587    479    158      38      62
##   F.Undergrad P.Undergrad Outstate Room.Board Books Personal PhD Terminal
## 1          2885          537    7440      3300    450    2200    70      78
## 2          2683          1227   12280      6450    750    1500    29      30
## 3          1036           99   11250      3750    400    1165    53      66
## 4           510           63   12960      5450    450     875    92      97
## 5           249          869    7560      4120    800    1500    76      72
## 6           678           41   13500      3335    500     675    67      73
##   S.F.Ratio perc.alumni Expend Grad.Rate
## 1        18.1         12    7041        60
## 2        12.2         16   10527        56
## 3        12.9         30    8735        54
## 4         7.7         37   19016        59
## 5        11.9          2   10922        15
## 6         9.4         11    9727        55
```

b. Next, I change the rownames to the university name and delete that column.

```
rownames(college) <- college[,1]
college <- college[,-1]
head(college)
```

```
##               Private Apps Accept Enroll Top10perc Top25perc
## Abilene Christian University   Yes 1660  1232   721      23      52
## Adelphi University            Yes 2186  1924   512      16      29
## Adrian College                Yes 1428  1097   336      22      50
## Agnes Scott College           Yes  417   349   137      60      89
## Alaska Pacific University      Yes  193   146    55      16      44
## Albertson College             Yes  587   479   158      38      62
##               F.Undergrad P.Undergrad Outstate Room.Board Books
## Abilene Christian University    2885         537   7440      3300  450
## Adelphi University             2683        1227  12280      6450  750
## Adrian College                 1036         99  11250      3750  400
## Agnes Scott College             510         63  12960      5450  450
## Alaska Pacific University        249        869   7560      4120  800
## Albertson College               678         41  13500      3335  500
##               Personal PhD Terminal S.F.Ratio perc.alumni Expend
## Abilene Christian University    2200   70      78    18.1      12  7041
## Adelphi University             1500   29      30    12.2      16 10527
## Adrian College                 1165   53      66    12.9      30  8735
## Agnes Scott College             875   92      97     7.7      37 19016
## Alaska Pacific University       1500   76      72    11.9       2 10922
## Albertson College              675   67      73     9.4      11  9727
##               Grad.Rate
## Abilene Christian University    60
## Adelphi University             56
## Adrian College                 54
## Agnes Scott College            59
## Alaska Pacific University       15
## Albertson College              55
```

c. Next we examine some stats on the data

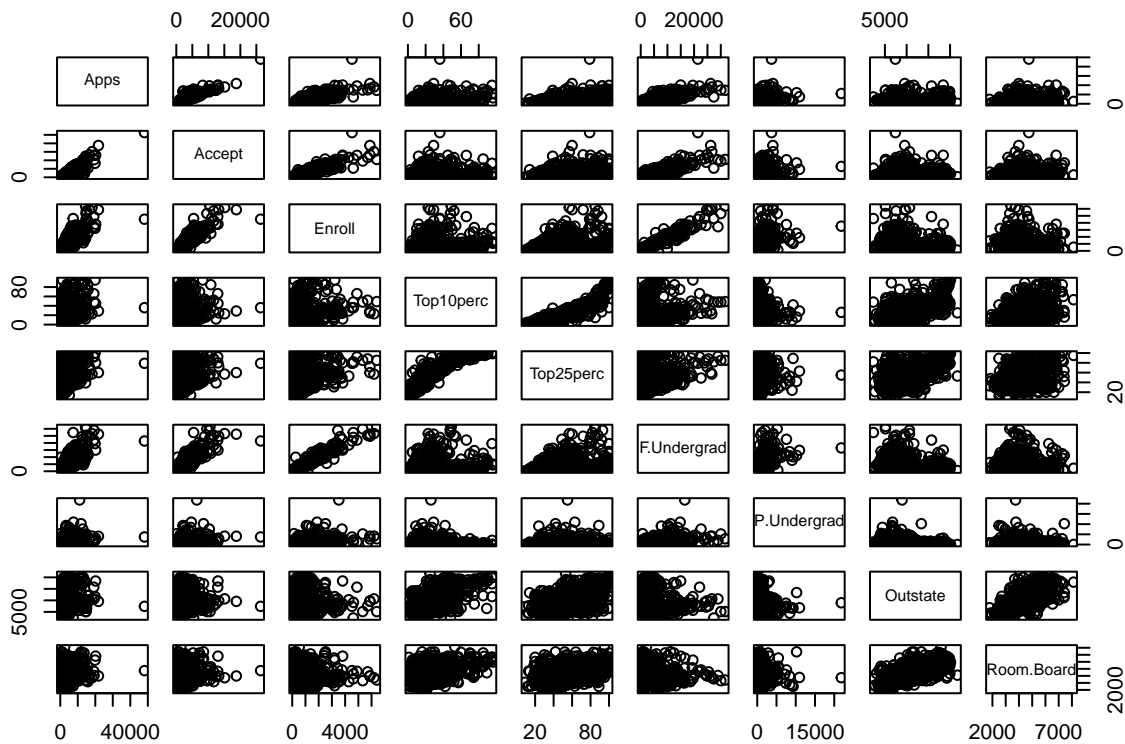
```
summary(college)
```

```
##      Private      Apps      Accept      Enroll
## Length:777      Min.   : 81      Min.   : 72      Min.   : 35
## Class :character 1st Qu.: 776      1st Qu.: 604      1st Qu.: 242
## Mode  :character Median : 1558      Median : 1110      Median : 434
##               Mean  : 3002      Mean  : 2019      Mean  : 780
##               3rd Qu.: 3624      3rd Qu.: 2424      3rd Qu.: 902
##               Max.   :48094      Max.   :26330      Max.   :6392
##      Top10perc      Top25perc      F.Undergrad      P.Undergrad
## Min.   : 1.00      Min.   : 9.0      Min.   : 139      Min.   : 1.0
## 1st Qu.:15.00      1st Qu.: 41.0      1st Qu.: 992      1st Qu.: 95.0
## Median :23.00      Median : 54.0      Median : 1707      Median : 353.0
## Mean   :27.56      Mean   : 55.8      Mean   : 3700      Mean   : 855.3
## 3rd Qu.:35.00      3rd Qu.: 69.0      3rd Qu.: 4005      3rd Qu.: 967.0
## Max.   :96.00      Max.   :100.0      Max.   :31643      Max.   :21836.0
##      Outstate      Room.Board      Books      Personal
## Min.   : 2340      Min.   :1780      Min.   : 96.0      Min.   : 250
## 1st Qu.: 7320      1st Qu.:3597      1st Qu.: 470.0      1st Qu.: 850
## Median : 9990      Median :4200      Median : 500.0      Median :1200
```

```
## Mean :10441 Mean :4358 Mean : 549.4 Mean :1341
## 3rd Qu.:12925 3rd Qu.:5050 3rd Qu.: 600.0 3rd Qu.:1700
## Max. :21700 Max. :8124 Max. :2340.0 Max. :6800
##      PhD      Terminal      S.F.Ratio      perc.alumni
## Min.   : 8.00   Min.   : 24.0   Min.   : 2.50   Min.   : 0.00
## 1st Qu.: 62.00   1st Qu.: 71.0   1st Qu.:11.50   1st Qu.:13.00
## Median : 75.00   Median : 82.0   Median :13.60   Median :21.00
## Mean   : 72.66   Mean   : 79.7   Mean   :14.09   Mean   :22.74
## 3rd Qu.: 85.00   3rd Qu.: 92.0   3rd Qu.:16.50   3rd Qu.:31.00
## Max.   :103.00   Max.   :100.0   Max.   :39.80   Max.   :64.00
##      Expend      Grad.Rate
## Min.   : 3186   Min.   : 10.00
## 1st Qu.: 6751   1st Qu.: 53.00
## Median : 8377   Median : 65.00
## Mean   : 9660   Mean   : 65.46
## 3rd Qu.:10830   3rd Qu.: 78.00
## Max.   :56233   Max.   :118.00
```

I am not familiar with the `pairs` command but here goes

```
pairs(college[,2:10])
```




That is a nice graphic although a bit too small for my tastes. I hope it compiles correctly in the pdf...

Next I'll create the boxplot for out of state tuition vs the public or private.

```
boxplot(Outstate ~ Private, data = college, main = "Out Of State Tuition by College Type", ylab = "Out of State Tuition")
```

## Out Of State Tuition by College Type



This looks fine although I do prefer ggplot2.

Next I comment the code as requested

```
Elite <- rep ("No", nrow(college )) #This creates a vector that full of No that is the same width as th
Elite [college$Top10perc >50] <- "Yes" #this changes some of the nos to yes if the top10 is more than 5
Elite <- as.factor (Elite) #this casts the vector as a factor vector. This is useful in that Elite now
college <- data.frame(college ,Elite) #this adds the column to the original dataframe and saves it

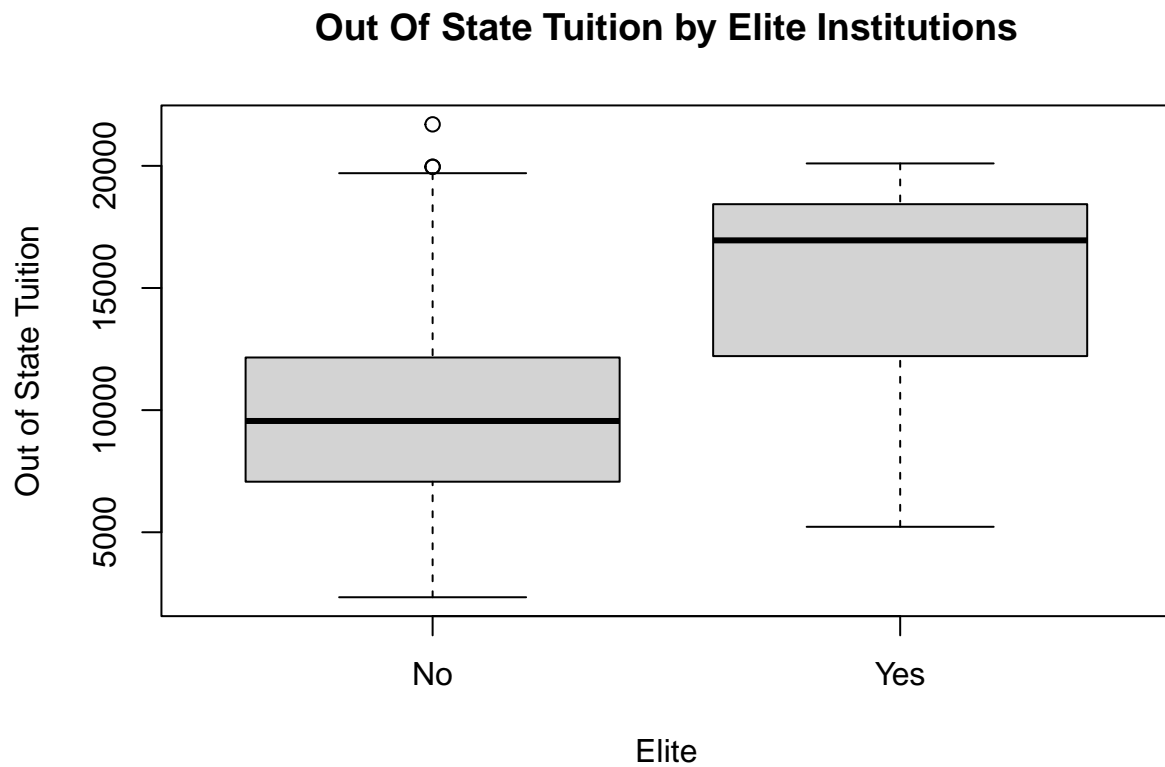
summary(Elite)
```

```
## No Yes
## 699 78
```

It appears that there are 78 elite universities. Let's explore tutions with this new factor

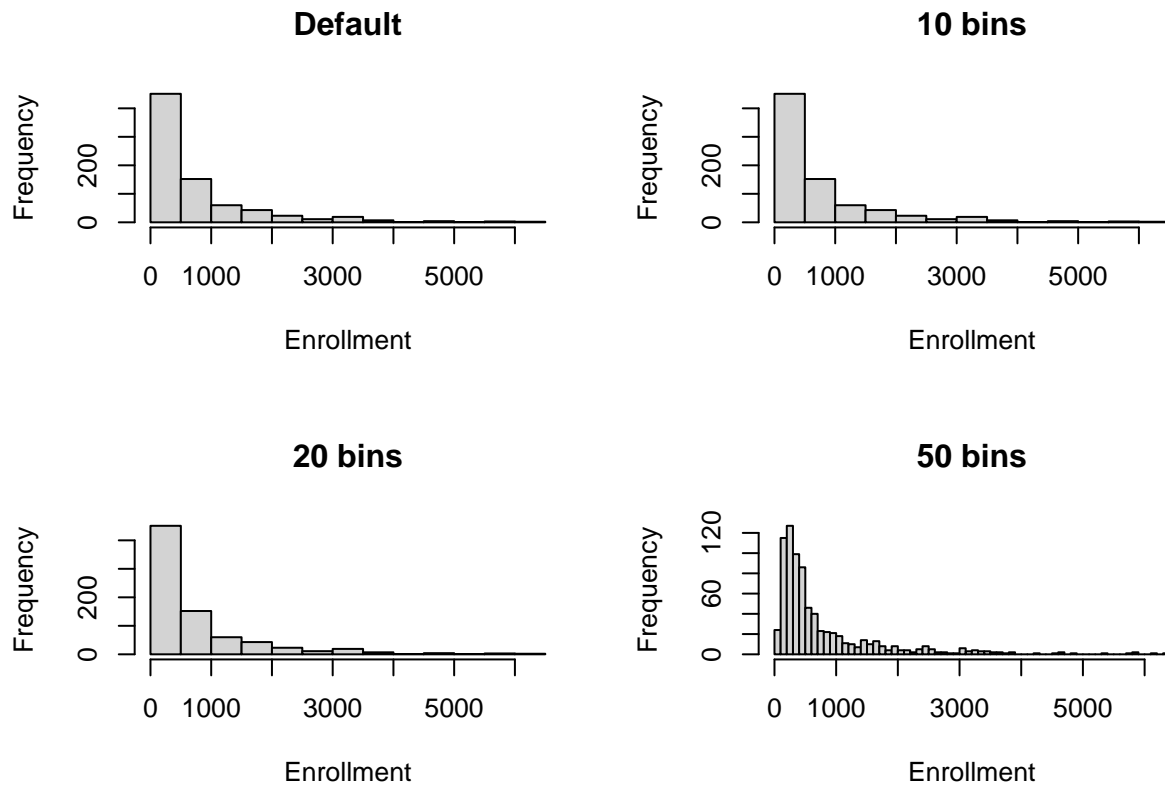
```
boxplot(Outstate ~ Elite, data = college, main = "Out Of State Tuition by Elite Institutions", ylab = "Out of State Tuition")
```





Next we look at a few histograms with differing number of bins.

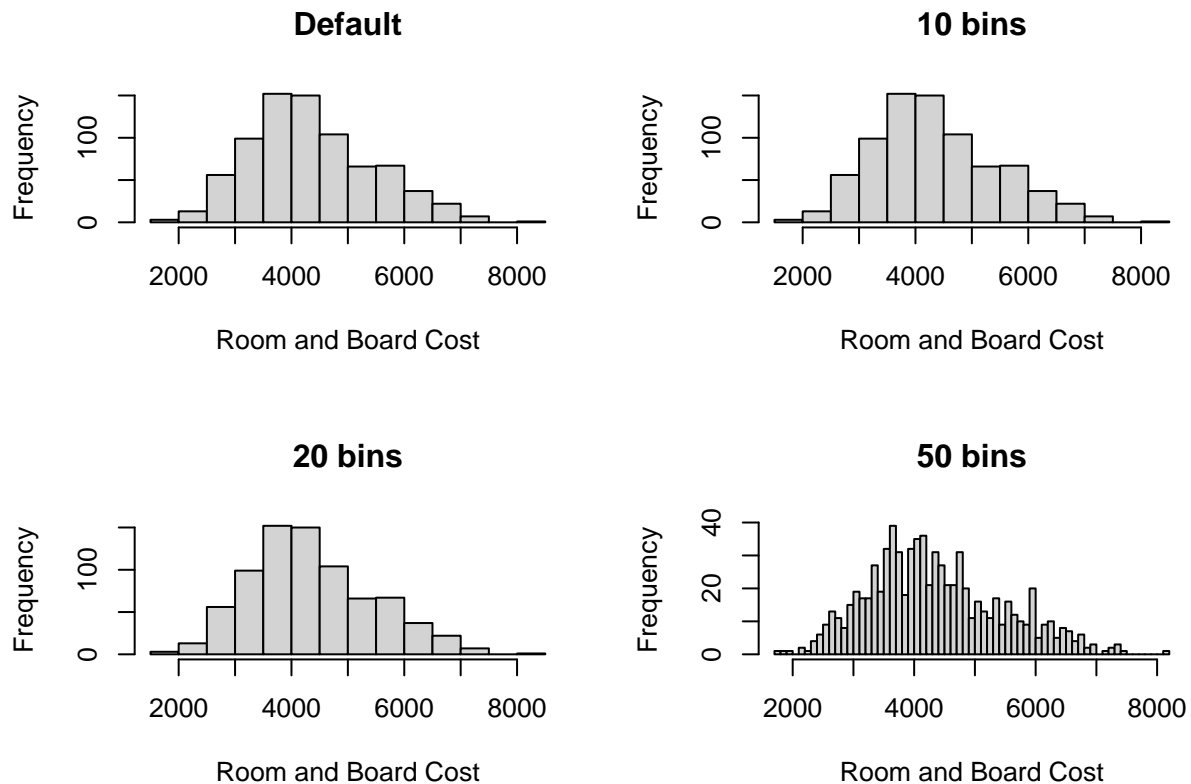
```
par(mfrow=c(2,2))
hist(college[, 'Enroll'], main = "Default", xlab = "Enrollment")
hist(college[, 'Enroll'], main = "10 bins", breaks = 10, xlab = "Enrollment")
hist(college[, 'Enroll'], main = "20 bins", breaks = 20, xlab = "Enrollment")
hist(college[, 'Enroll'], main = "50 bins", breaks = 50, xlab = "Enrollment")
```



I don't see much difference between the default, 10 nor 20. The 50 does look a bit different.

Again just to try it once more

```
par(mfrow=c(2,2))
hist(college[, 'Room.Board'], main = "Default", xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "10 bins", breaks = 10, xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "20 bins", breaks = 20, xlab = "Room and Board Cost")
hist(college[, 'Room.Board'], main = "50 bins", breaks = 50, xlab = "Room and Board Cost")
```



It kind of looks like more than 10 breaks. Maybe the default overrides that option if you set it too low...

## 4 Using R: Manipulating Data in Data Frames

- a. First, I'll load some data directly from a package. This `baseball` data comes from the `plyr` package loaded earlier.

```
head(baseball)
```

```
##           id year stint team lg  g  ab  r  h X2b X3b hr rbi sb cs bb so ibb
## 4   ansonca01 1871     1  RC1   25 120 29 39  11   3  0  16  6  2  2  1  NA
## 44  forceda01 1871     1  WS3   32 162 45 45   9   4  0  29  8  0  4  0  NA
## 68  mathebo01 1871     1  FW1   19  89 15 24   3   1  0  10  2  1  2  0  NA
## 99  startjo01 1871     1  NY2   33 161 35 58   5   1  1  34  4  2  3  0  NA
## 102 suttoez01 1871     1  CL1   29 128 35 45   3   7  3  23  3  1  1  0  NA
## 106 whitede01 1871     1  CL1   29 146 40 47   6   5  1  21  2  2  4  1  NA
##      hbp sh sf gidp
## 4      NA NA NA   NA
## 44     NA NA NA   NA
## 68     NA NA NA   NA
## 99     NA NA NA   NA
## 102    NA NA NA   NA
## 106    NA NA NA   NA
```

- b. Lots of baseball data!

```
baseball[baseball$year<1954,'sf'] = 0 #set all sf before 1954 to 0
baseball[is.na(baseball$hbp),'hbp'] = 0 #set all null values for hit by pitch to 0
```

```
baseball <- baseball[baseball$ab>=50,]
```

c. Now that the data is clean, we will apply the obp formula of

$$obp = \frac{h + bb + hbp}{ab + bb + hbp + sf}$$

```
baseball <- mutate(baseball, obp = (h+bb+hbp)/(ab+bb+hbp+sf))
```

```
head(baseball)
```

```
##           id year stint team lg  g  ab  r  h X2b X3b hr rbi sb cs bb so ibb
## 4  ansonca01 1871     1  RC1   25 120 29 39  11   3  0  16  6  2  2  1  NA
## 44 forceda01 1871     1  WS3   32 162 45 45   9   4  0  29  8  0  4  0  NA
## 68 mathebo01 1871     1  FW1   19  89 15 24   3   1  0  10  2  1  2  0  NA
## 99 startjo01 1871     1  NY2   33 161 35 58   5   1  1  34  4  2  3  0  NA
## 102 suttoez01 1871     1  CL1   29 128 35 45   3   7  3  23  3  1  1  0  NA
## 106 whitede01 1871     1  CL1   29 146 40 47   6   5  1  21  2  2  4  1  NA
##           hbp sh sf gidp           obp
## 4           0 NA  0    NA 0.3360656
## 44          0 NA  0    NA 0.2951807
## 68          0 NA  0    NA 0.2857143
## 99          0 NA  0    NA 0.3719512
## 102         0 NA  0    NA 0.3565891
## 106         0 NA  0    NA 0.3400000
```

d. Now that we have that info added, let's find the top five players for obp of all time.

```
arrange(baseball, -obp)[1:5,c('year','id','obp')] #I get the top records with 1:5, restrict on to the c
```

```
##   year      id      obp
## 1 2004 bondsba01 0.6094003
## 2 2002 bondsba01 0.5816993
## 3 1941 willite01 0.5528053
## 4 1899 mcgrajo01 0.5474860
## 5 1923 ruthba01 0.5445402
```

We see here Barry Bonds (from the 'roids era twice), Ted Williams(a year he hit .400), John McGraw (a player I was not familiar with though he did have a season with my home team Cardinals in 1900) and the babe himself Babe Ruth.

## 5 Using R: aggregate() Function


a. I am going to grab the quakes dataset.

```
head(quakes)
```

```
##      lat    long depth mag stations
## 1 -20.42 181.62   562 4.8         41
## 2 -20.62 181.03   650 4.2         15
## 3 -26.00 184.10    42 5.4         43
## 4 -17.97 181.66   626 4.1         19
## 5 -20.42 181.96   649 4.0         11
## 6 -19.68 184.31   195 4.0         12
```

b. Next we will examine magnitude versus depth with a scatter plot.

```
plot(quakes$depth,quakes$mag, xlab = 'Depth', ylab = 'Magnitude', main = 'Scatter Plot of Depth vs Magnitude')
```



c. Next we will aggregate the data to look at the average depth for each of the magnitude levels

```
quakeAvgDepth = aggregate(quakes$depth, list(mag = quakes$mag), mean)
```

Not too bad when you follow the example in the help menu.

d. Next I rename the dataframe to have useful column names and print it to see the nice output.


```
colnames(quakeAvgDepth) = c('mag','meanDepth')
head(quakeAvgDepth)
```

```
##   mag meanDepth
## 1 4.0   410.0652
## 2 4.1   412.4000
## 3 4.2   389.8778
## 4 4.3   357.9294
## 5 4.4   307.1188
## 6 4.5   333.6729
```

e. Now we plot again to see if there is a relationship in the aggregate

```
plot(quakeAvgDepth$mag,quakeAvgDepth$meanDepth, xlab = 'Magnitude', ylab = 'Average Depth',main = 'Scat
```

### Scatter of Aggregated Magnitude vs Mean Depth



- f. There clearly appears to be a relationship here. It was not as obvious in the full data case but the relationship appears in the aggregate. I do question a bit of this methodology though. We are aggregating a continuous variable that has been truncated to two decimals. Richter scale (magnitude) is a famous example of a logarithmic scale so small rounding errors are amplified in varying degrees as you increase the scale. While yes, I believe there is a relationship, I'd be worried about generalizing too far based on this data.